

Wampa World

This Wampa World logic-based agent homework was developed by [Joseph Willem Ricci](#) and [Ben Swanson](#) upon [the original Jupyter Notebook](#) by [Lara Martin](#).

Gameplay

R2-D2 begins at the bottom-left location $(0, 0)$ and must navigate the rectangular grid of unknown size. To navigate the grid, **he** can turn **left** or **right**, or move **forward**. R2-D2's goal is to rescue Luke by navigating to the room that contains Luke, **grabbing** him, and navigating back to $(0, 0)$ to **climb** out of the cave.

Along the way, R2-D2 must avoid pits which he can fall into, and must avoid the Wampa, which can destroy him. In any adjacent room to a pit is a **breeze**. In any adjacent room to a Wampa is a **stench**. There can be $[0, 1]$ Wampas, and there can be $[0, m \times n - 2]$ pits, where the grid is of size $m \times n$. Each room can have 0 or 1 features from **[luke, pit, wall, wampa]**.

R2-D2 is also carrying a blaster with one shot and infinite range, and can **shoot** the Wampa with a shot in its direction. If the Wampa is killed by the shot, a **scream** can be perceived in every room.

A **gasp** from Luke can be perceived by R2-D2 if they are both in the same room.

Finally, if R2-D2 moves **forward** into a wall, R2-D2 perceives a **bump**. For this assignment, assume that R2 knows that the world is rectangular and that there are no internal walls. A hypothetical 1x1 grid would have walls at

```
{(-1, -1), (-1, 0), (-1, 1), (0, -1), (0, 1), (1, -1), (1, 0), (1, 1)}
```

Percepts:

```
['stench', 'breeze', 'gasp', 'bump', 'scream']
```

Where at some location, R2's percepts might look like

```
[None, 'breeze', None, 'bump', None]
```

Actions:

left, right, forward, grab, climb, shoot

To run the game, run the following command in your terminal from your homework's directory:

```
`python wampa_world.py <scenario>`
```

File Structure

scenarios.py

Contains five scenarios, S1, S2, S3, S4 and S5 to test your program with. Feel free to write your own!

utils.py

Contains miscellaneous helper and utility functions. You may want to utilize `flatten(tup)` and `get_direction(degrees)`.

visualize__world.py

Is called during gameplay to visualize the current state of the world.

wampa__world.py

Contains the WampaWorld class which defines gameplay and the main gameplay loop.

agent.py

This is where you'll be implementing your functions. Contains R2-D2's constructor including initial knowledge base class, KB. Familiarize yourself with the knowledge base, as the agent class will be using it for logical inference.

agent.py TODOs:

`adjacent_rooms(self, room):`

Returns a set of tuples representing all adjacent rooms to 'room'.

Input: `(2, 3)`

Expected output: `{(1, 3), (3, 3), (2, 2), (2, 4)}`

`record_percepts(self, sensed_percepts, current_location):`

Update the percepts in agent's KB with the percepts sensed in the current location, and update visited_rooms and all_rooms accordingly.

`enumerate_possible_worlds(self):`

Return the set of all possible worlds, where a possible world is a tuple of `(pit_rooms, wampa_room)`, `pit_rooms` is a tuple of tuples representing possible pit rooms, and `wampa_room` is a tuple representing a possible wampa room. Since the goal is to combinatorially enumerate all the possible worlds (pit and wampa locations) over the set of rooms that could potentially have a pit or a wampa, we first want to find that set. To do that, subtract the set of rooms that you know cannot have a pit or wampa from the set of all rooms. For example, you know that a room with a wall cannot have a pit or wampa. Then use `itertools.combinations` to return the set of possible worlds, or all combinations of possible pit and wampa locations. You may find the `utils.flatten(tup)` method useful here for flattening `wampa_room` from a tuple of tuples into a tuple. The output of this function will be queried to find the model of the query, and will be checked for consistency with the KB to find the model of the KB.

From step 0, the initial position on scenario S1:

```
. . . P
W L P .
. . . .
^ . P .
```

Expected output `{(), ()}`. Since all adjacent rooms are known to be safe when there is no breeze and no stench in the current room, so all rooms in `KB.all_rooms` are known to be safe, so there are no possible worlds with pits or wampas.

From step 1, the following position resulting from a forward action:

```
. . . P
W L P .
^ . . .
. . P .
```

Expected output, `possible_worlds =`

```
{
    ((), ()),
    ((), (-1, 1)),
    ((), (0, 2)),
    ((), (1, 1)),
    (((-1, 1),), ()),
    (((-1, 1),), (0, 2)),
    (((-1, 1),), (1, 1)),
    (((-1, 1), (0, 2)), ()),
    (((-1, 1), (0, 2)), (1, 1)),
    (((-1, 1), (1, 1)), ()),
    (((-1, 1), (1, 1)), (0, 2)),
    (((-1, 1), (1, 1), (0, 2)), ()),
    (((0, 2),), ()),
    (((0, 2),), (-1, 1)),
    (((0, 2),), (1, 1)),
    (((1, 1),), ()),
    (((1, 1),), (-1, 1)),
    (((1, 1),), (0, 2)),
    (((1, 1), (0, 2)), ()),
    (((1, 1), (0, 2)), (-1, 1))
}
```

which corresponds to every combination of pit locations and wampa location in the rooms that could potentially have a pit or wampa, `(-1, 1)`, `(0, 2)`, `(1, 1)`. Note, this is not doing “model checking”. This is simply enumerating all possible pit and wampa locations. This set will be used to check against our KB and to query with our queries.

`pit_room_is_consistent_with_KB(self, room):`

Return `True` if the room could be a pit given KB, `False` otherwise. A room could be a pit if all adjacent rooms that have been visited have breeze. A room cannot be a pit if any adjacent rooms that have been visited have not had breeze perceived in them. This will be used to find the model of the KB.

`wampa_room_is_consistent_with_KB(self, room):`

Return `True` if the room could be a wampa given KB, `False` otherwise. A room could be a wampa if all adjacent rooms that have been visited have stench. A room cannot be a wampa if any adjacent rooms that have been visited have not had stench perceived in them. This will be used to find the

model of the KB.

```
find_model_of_KB(self, possible_worlds):
```

Return the subset of all possible worlds consistent with KB. `possible_worlds` is a set of tuples `(pit_rooms, wampa_room)`, `pit_rooms` is a set of tuples of possible pit rooms, and `wampa_room` is a tuple representing a possible wampa room. A world is consistent with the KB if `wampa_room` is consistent and all pit rooms are consistent with the KB.

Input: `possible_worlds`, from step 1.

Output, `model_of_KB` =

```
{
    (( ), (-1, 1)),
    (( ), (0, 2)),
    (( ), (1, 1))
}
```

which is the subset of worlds from `possible_worlds` that “checked out” with the model of the KB. No breeze has been perceived yet, so the only worlds that are consistent with the KB are those with no pits. And with stench perceived in (0, 1), rooms `{(-1, 1), (0, 2), (1, 1)}` are the only rooms for which `wampa_room_is_consistent_with_KB` returns `True`.

If we consider the state of S1 after taking the following actions from the initial position: `forward, left, left, forward, left, forward`

To end up in the following state:

```
. . . P
W L P .
. . . .
. > P .
```

After recording percepts in that position, we would expect `model_of_KB` =

```
{
    (((1, -1),), (-1, 1)),
    (((1, -1),), (0, 2)),
    (((2, 0),), (-1, 1)),
    (((2, 0),), (0, 2)),
    (((2, 0), (1, -1)), (-1, 1)),
    (((2, 0), (1, -1)), (0, 2))
}
```

which is all of the possible worlds `(pit_rooms, wampa_room)` where `pit_room_is_consistent_with_KB(room)` returns `True` for all rooms in `pit_rooms` and `wampa_room_is_consistent_with_KB(room)` returns `True` for `wampa_room`.

```
find_model_of_query(self, query, room, possible_worlds):
```

Where query can be “pit_in_room”, “wampa_in_room”, “no_pit_in_room” or “no_wampa_in_room”, filter the set of worlds according to the query and room.

Inputs: `"wampa_in_room"`, `(0, 2)`, `possible_worlds` from step 1.

Output:

```
{
    (((), (0, 2)),
    (((-1, 1),), (0, 2)),
    (((-1, 1), (1, 1)), (0, 2)),
    (((1, 1),), (0, 2))
}
```

which is the subset of `possible_worlds` that contain `(0, 2)` in the `wampa_room` index.

```
inference_algorithm(self):
```

First, make some basic inferences:

1. If there is no breeze or stench in current location, infer that the adjacent rooms are safe.
2. Infer wall locations given bump percept.
3. Infer Luke's location given gasp percept.
4. Infer whether the Wampa is alive given scream percept.

Then, infer whether each adjacent room is safe, pit or wampa by following the backward-chaining resolution algorithm:

1. Enumerate possible worlds.
2. Find the model of the KB, i.e. the subset of possible worlds consistent with the KB.
3. For each adjacent room and each query, find the model of the query.
4. If the model of the KB is a subset of the model of the query, the query is entailed by the KB.
5. Update `KB.pits`, `KB.wampa`, and `KB.safe_rooms` based on any newly derived knowledge.

```
all_safe_next_actions(self)
```

Define R2-D2's valid and safe next actions based on his current location and knowledge of the environment.

```
choose_next_action(self)
```

Choose next action from all safe next actions. You may want to prioritize some actions based on current state. For example, if R2-D2 knows Luke's location and is in the same room as Luke, you may want to prioritize `grab` over all other actions. Similarly, if R2-D2 has Luke, you may want to prioritize moving toward the exit. You can implement this as basically (randomly choosing between safe actions) or as sophisticated (optimizing exploration of unvisited states, finding shortest paths, etc.) as you like.

Submission

Upload your file `agent.py` to the Gradescope assignment submission.

A (non-exhaustive) suite of test cases for each method can help you troubleshoot and ascertain whether you are along the right track.